UDACITY

‹ Return to "Deep Reinforcement Learning Nanodegree" in the classroom

DISCUSS ON STUDENT HUB ›

Collaboration and Competition

REVIEW

CODE REVIEW

HISTORY

Requires Changes

1 SPECIFICATION REQUIRES CHANGES

You have done an exceptional job on this project.
Just need to fulfil some minor requirements that I mentioned.
I hope you keep your spirits up and keep working like this.
All the BEST 😄

Training Code

✓

The repository includes functional, well-documented, and organized code for training the agent.

All the code files are included in the repository

✓

The code is written in PyTorch and Python 3.

Python3 and PyTorch have been used.

✓

The submission includes the saved model weights of the successful agent.

Saved model weights are included.

README

The GitHub submission includes a `README.md` file in the root of the repository.

README is included in the repository.

The README describes the the project environment details (i.e., the state and action spaces, and when the environment is considered solved).

README is informative, all the details about the environment are covered

The README has instructions for installing dependencies or downloading needed files.

Instructions to install the dependencies and downloading the required files have been provided.

The README describes how to run the code in the repository, to train the agent. For additional resources on creating READMEs or using Markdown, see here and here.

Report

The submission includes a file in the root of the GitHub repository (one of `Report.md`, `Report.ipynb`, or `Report.pdf`) that provides a description of the implementation.

The report is included in the repository.

The report clearly describes the learning algorithm, along with the chosen hyperparameters. It also describes the model architectures for any neural networks.

You also need to describe the learning algorithm in the report file.

A plot of rewards per episode is included to illustrate that the agents get an average score of +0.5 (over 100 consecutive episodes, after taking the maximum over both agents).

The submission reports the number of episodes needed to solve the environment.

Great work in achieving an average score of +0.5

✓

The submission has concrete future ideas for improving the agent's performance.

Great ideas for future improvement. Few more future ideas to explore:

- Use parameter space noise rather than noise on action. https://vimeo.com/252185862https://github.com/jvmancuso/ParamNoise
- We can use prioritised experience buffer. https://github.com/Damcy/prioritized-experience-replay
- Different replay buffer for actor/critic
- Try adding dropouts in critic network
- Turn off OU noise and use random noise

RESUBMIT

DOWNLOAD PROJECT

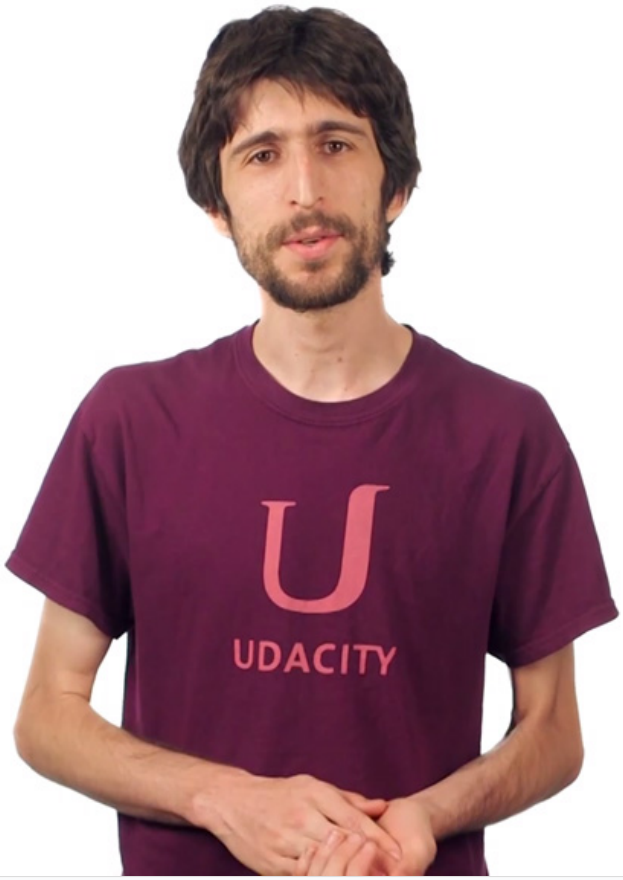

Best practices for your project resubmission

Ben shares 5 helpful tips to get you through revising and resubmitting your project.

Watch Video (3:01)

RETURN TO PATH